

# CoolPIM: Thermal-Aware Source Throttling for Efficient PIM Instruction Offloading

Lifeng Nai<sup>\*†</sup> Ramyad Hadidi<sup>†</sup> He Xiao<sup>†</sup> Hyojong Kim<sup>†</sup>  
Jaewoong Sim<sup>‡</sup> Hyesoon Kim<sup>†</sup>

<sup>†</sup>Georgia Institute of Technology

\* Google, <sup>‡</sup>Intel Labs

# Processing-in-Memory

2/24

Processing-in-memory (PIM) is regaining attention for energy efficient computing

- Graph Workloads: Data-Intensive, Little Data Reuse

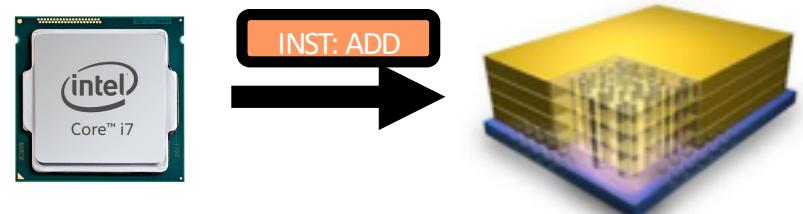


## Basic Concept: Offload compute to memory

- Reduce *costly* energy consumption of data movement
- Enable using large internal memory bandwidth



Conventional Data Processing



Processing-in-Memory

# Thermal Challenge in PIM

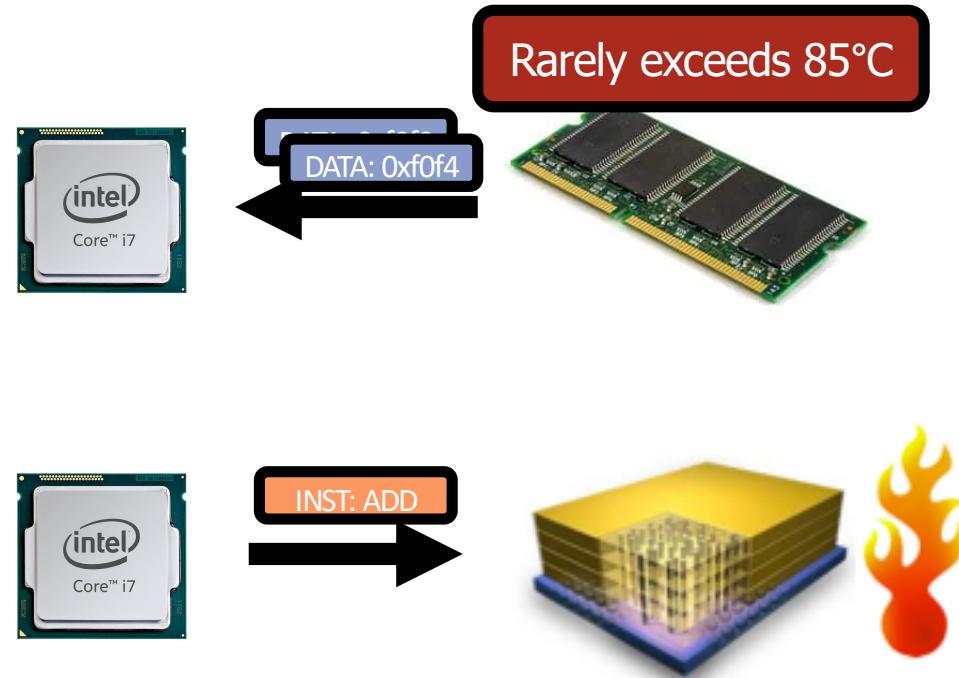
3/24

PIM could increase memory temperature beyond normal operating temperature (85°C)

- High BW (hundreds of GBs ~ TBs) from 3D-stacked memory
- Less effective heat transfer compared to DIMMs
- **PIM would make these thermal problems worse!**

Too Hot Memory Stack?

- Slower processing for memory requests
- Decreasing overall system performance



CoolPIM keeps the memory “*Cool*”  
to achieve better PIM performance

# Outline

---

4/24

## Introduction

## Hybrid Memory Cube

- Background
- Thermal Measurements & Thermal Modeling of Future HMC

## CoolPIM

- Software-Based Throttling
- Hardware-Based Throttling

## Evaluation

## Conclusion

# Background

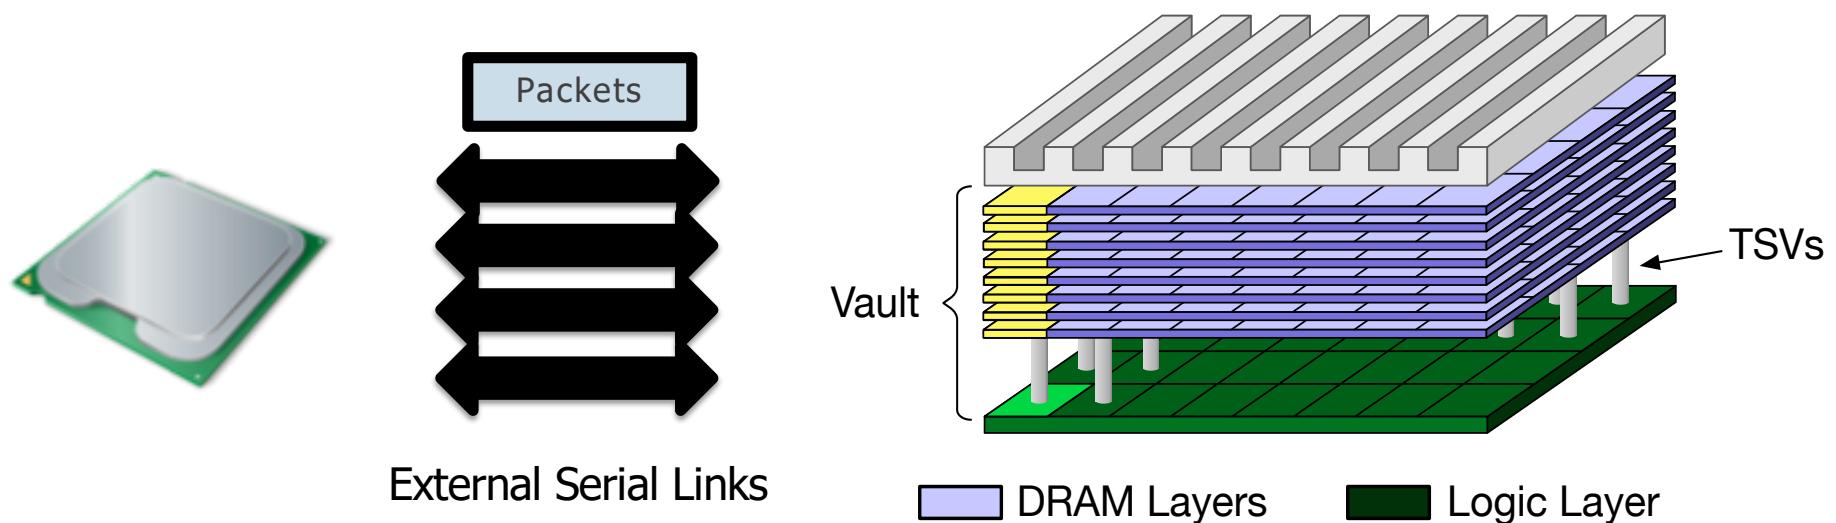
## Hybrid Memory Cube (HMC)

5/24

### A Hybrid Memory Cube (HMC) from Micron

- Multiple 3D-stacked DRAM layers + one logic layer with TSVs
- Vaults: equivalent to memory channels
- Full-duplex serial links between the host and HMC

No PIM functionality for existing HMC products yet



# Background

## PIM Offloading in Future HMC

6/24

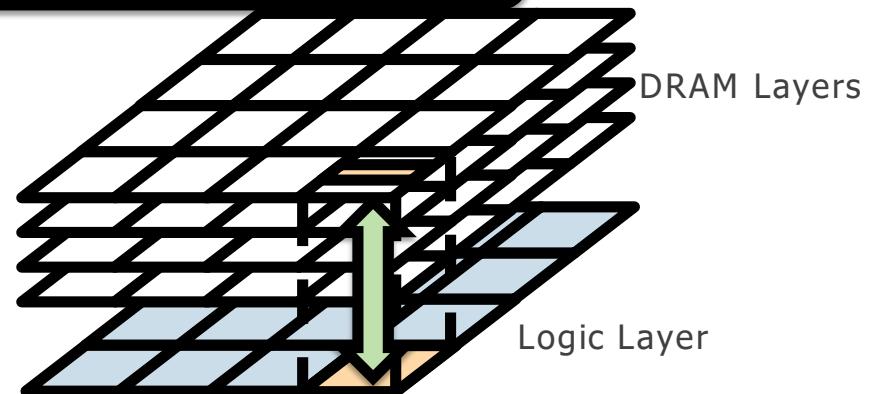
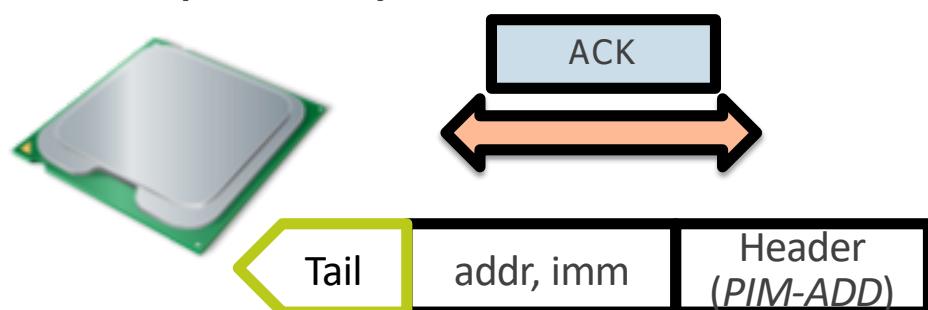
### Instruction-level PIM supported in future HMC (HMC 2.0)

- Perform **Read-Modify-Write (RMW)** operations atomically
- Similar to READ/WRITE packets; just different **CMD** in the **Header**
- **No HMC 2.0 product yet!**

Type	HMC 2.0 PIM Instruction
Arithmetic	Signed Add
Bitwise	Swap, bit write
Boolean	AND/NAND/OR/NOR/XOR

**Q: Can we offload all the PIM operations to HMC?  
What is the thermal impact of PIM in future HMC?**

*PIM-ADD (addr, imm)*



# Existing HMC Thermal Measurement (1)

7/24

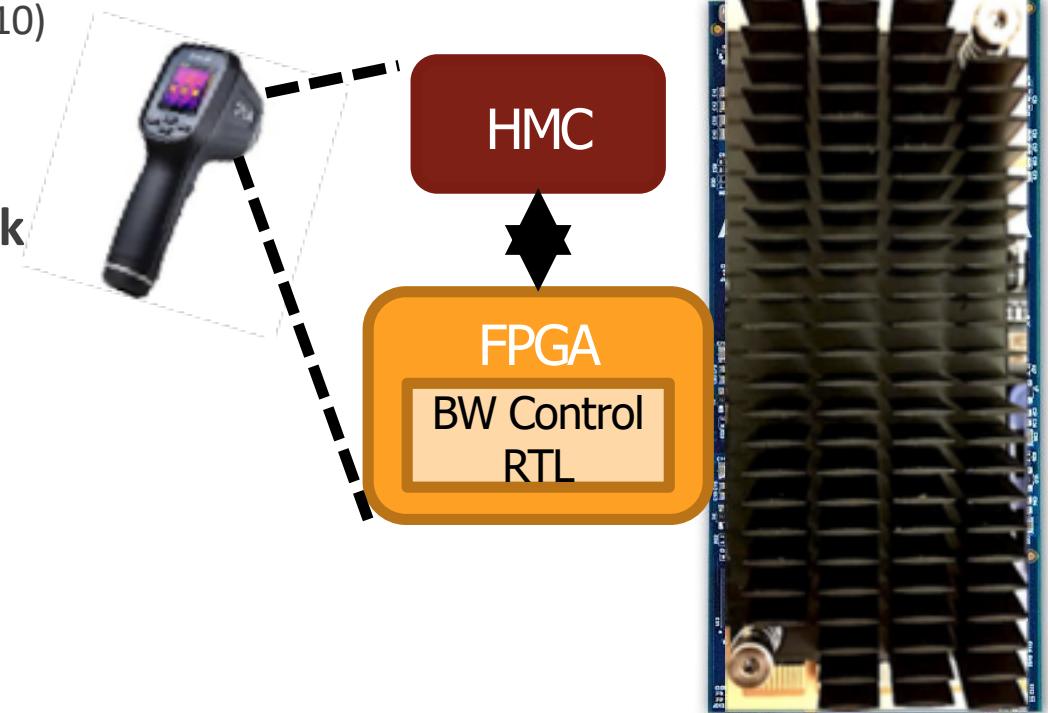
## Experiment Platform (Pico SC-6 Mini System)

- Intel Core i7 + FPGA Compute Modules (AC-510)
  - ▶ AC-510: 4GB HMC 1.1, Kintex Ultrascale

## Measure the temperature on the heat sink

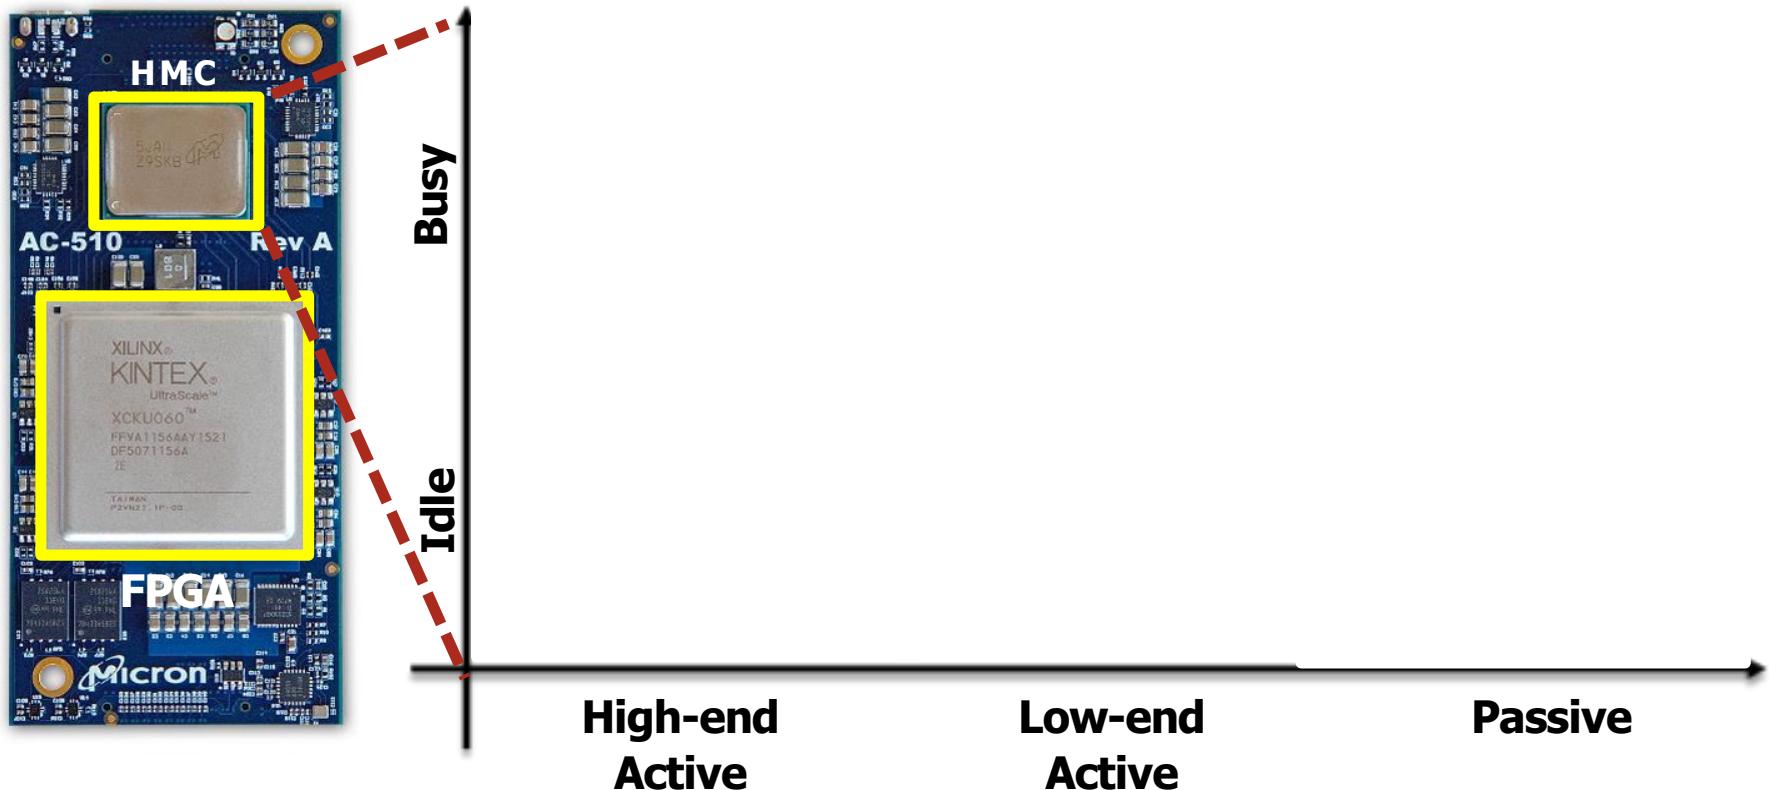
- Controlling memory BW via FPGA
- Applying three different cooling methods
  - ▶ High-End Active Heat Sink
  - ▶ Low-End Active Heat Sink
  - ▶ Passive Heat Sink

HMC 1.1 has no PIM functionality!



# Existing HMC Thermal Measurement (2)

8/24

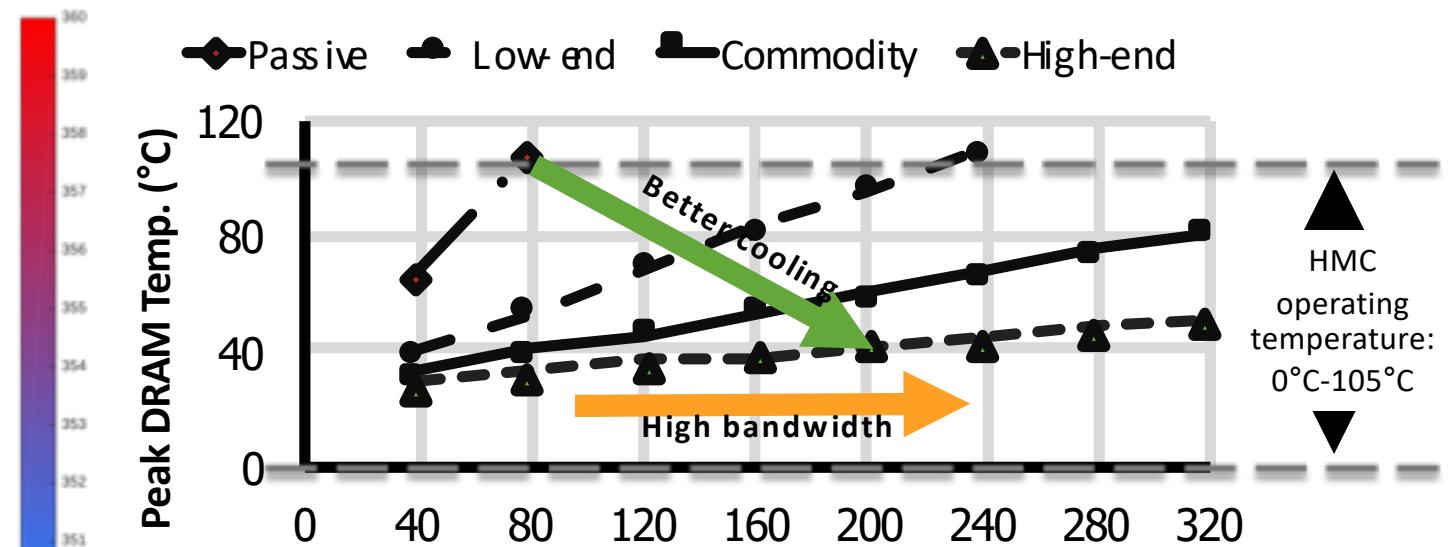
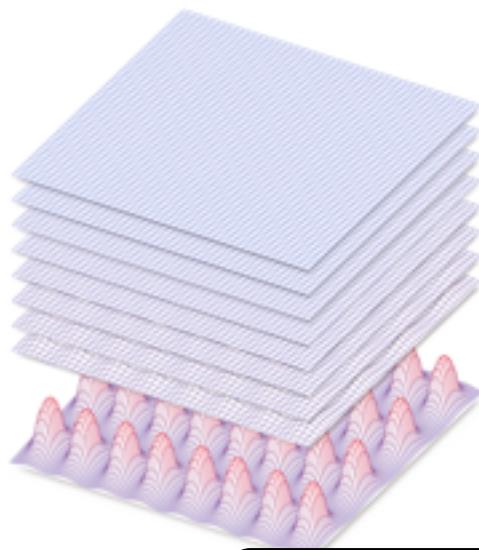


# Future HMC Thermal Modeling

9/24

## Thermal modeling for HMC 2.0 with commodity-server active cooling

- HMC 2.0 (w/o PIM) would reach 81°C at a full external BW (320GB/s)
  - We validated our thermal model against the measurements on HMC 1.1



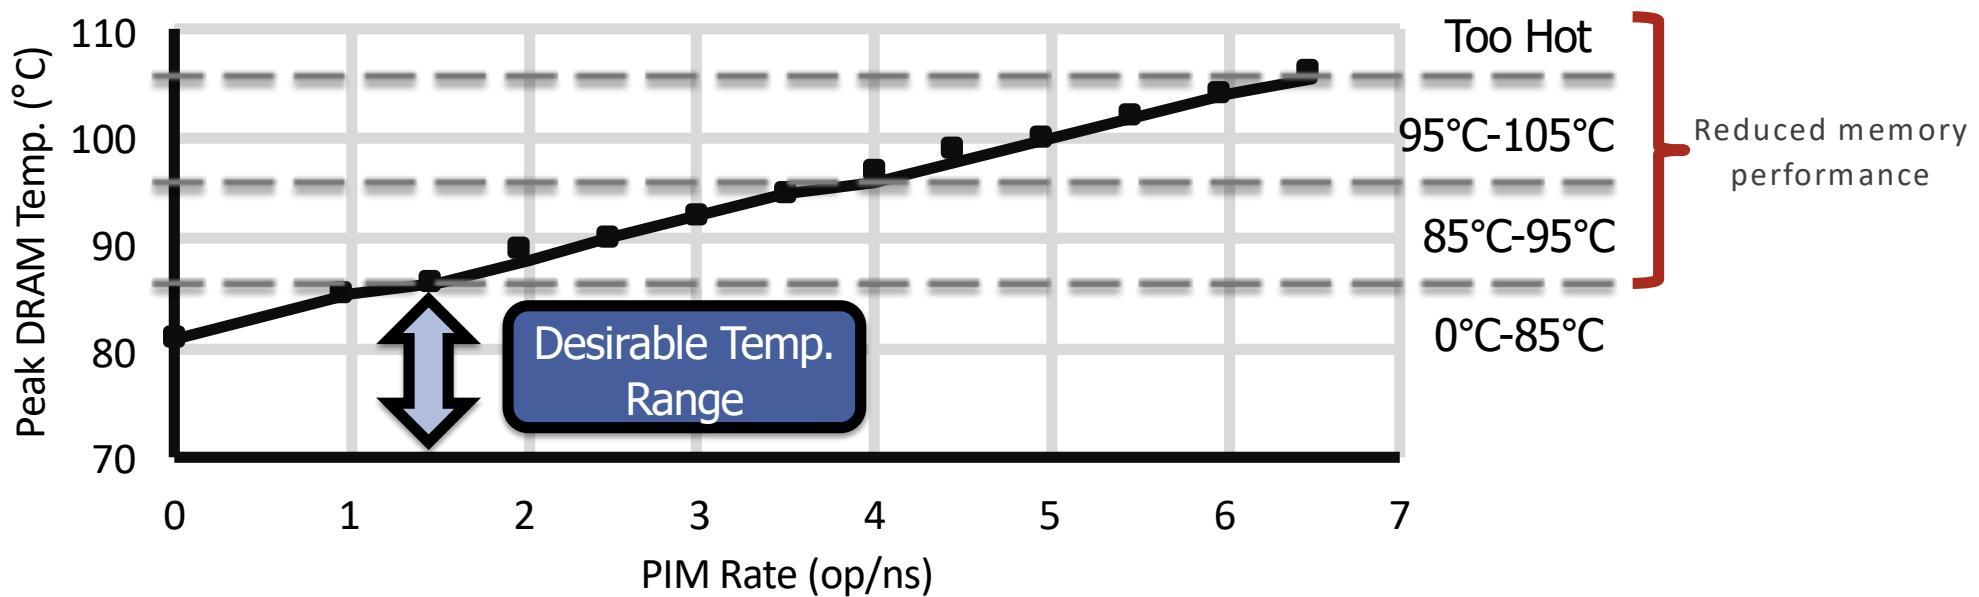
We need at least commodity-server cooling to benefit from PIM!

# Thermal Impact with PIM in HMC 2.0

10/24

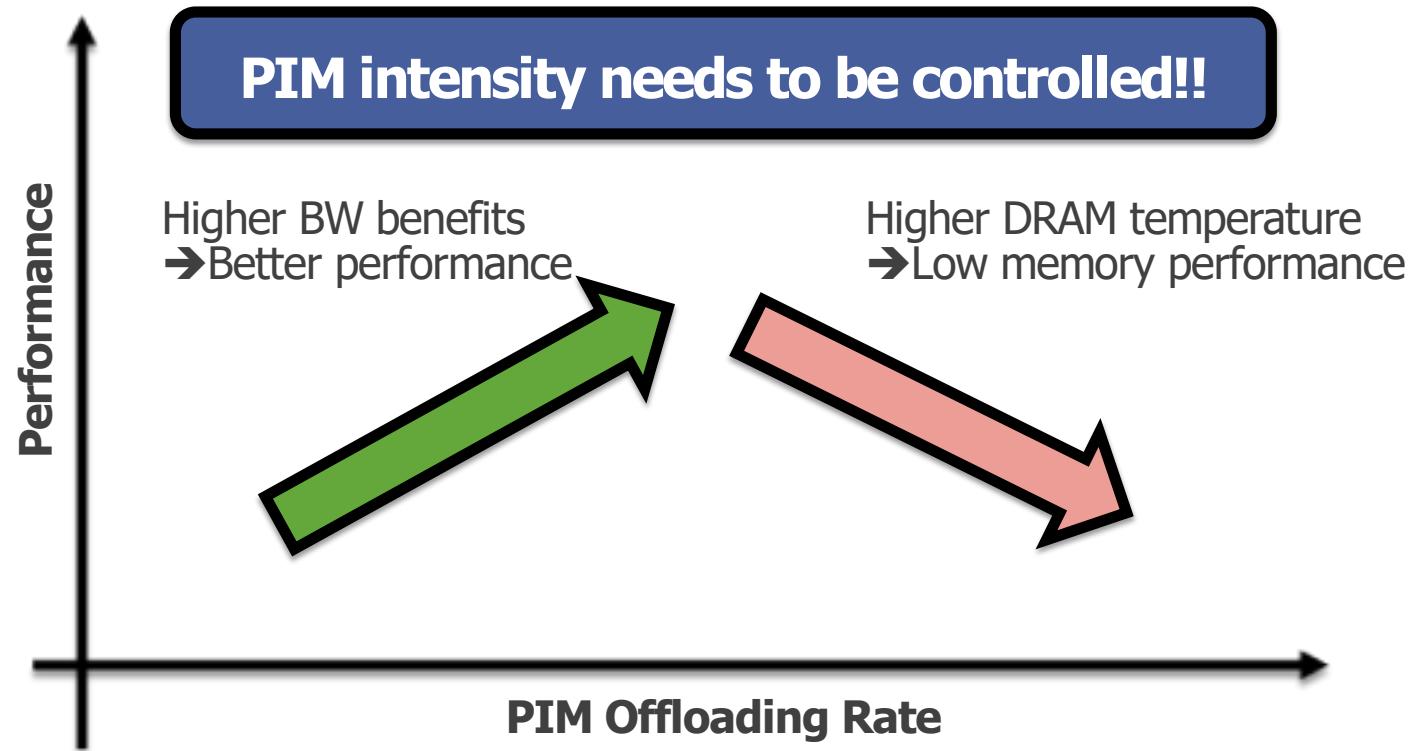
PIM increases memory temperature due to power consumption of logic and DRAM layers.

- In our modeling, the maximum PIM offloading rate is 6.5 PIM ops/ns
- A high offloading rate could reduce memory performance for cool down



# Performance Trade-off of PIM

11/24



## CoolPIM

**Controls PIM Intensity with Thermal Consideration**

# CoolPIM: Overview

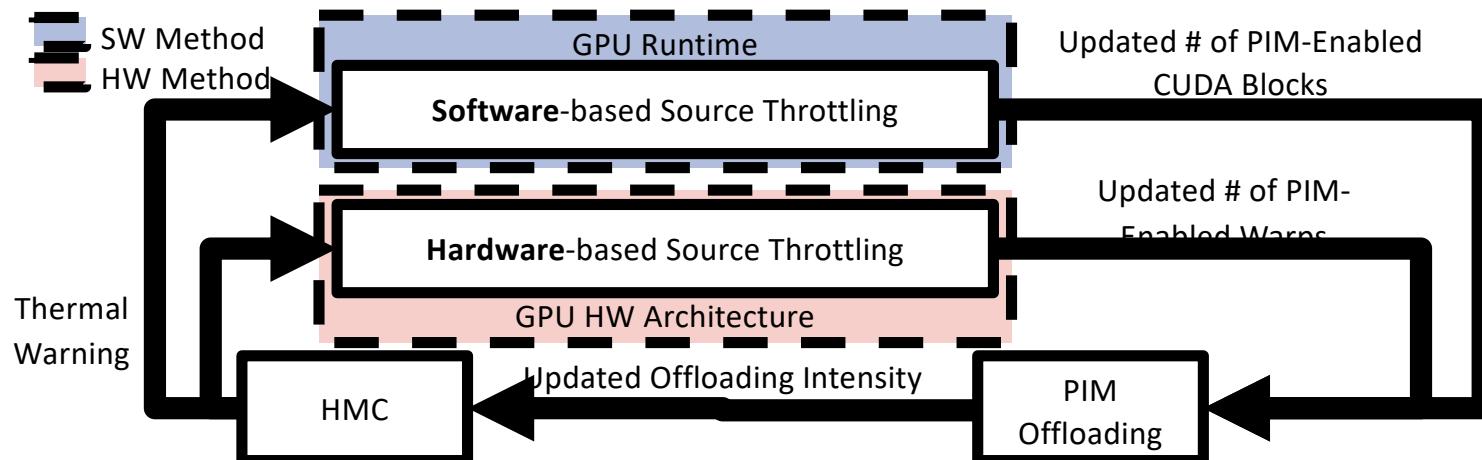
13/24

We propose two methods for GPU/HMC

- 1) A **SW** mechanism with **no hardware changes**
- 2) A **HW** mechanism with **changes in GPU architectures**

## Dynamic source throttling based on thermal warning messages from HMC

- Thermal warning -> lowers PIM intensity -> reduces internal temperature of HMC

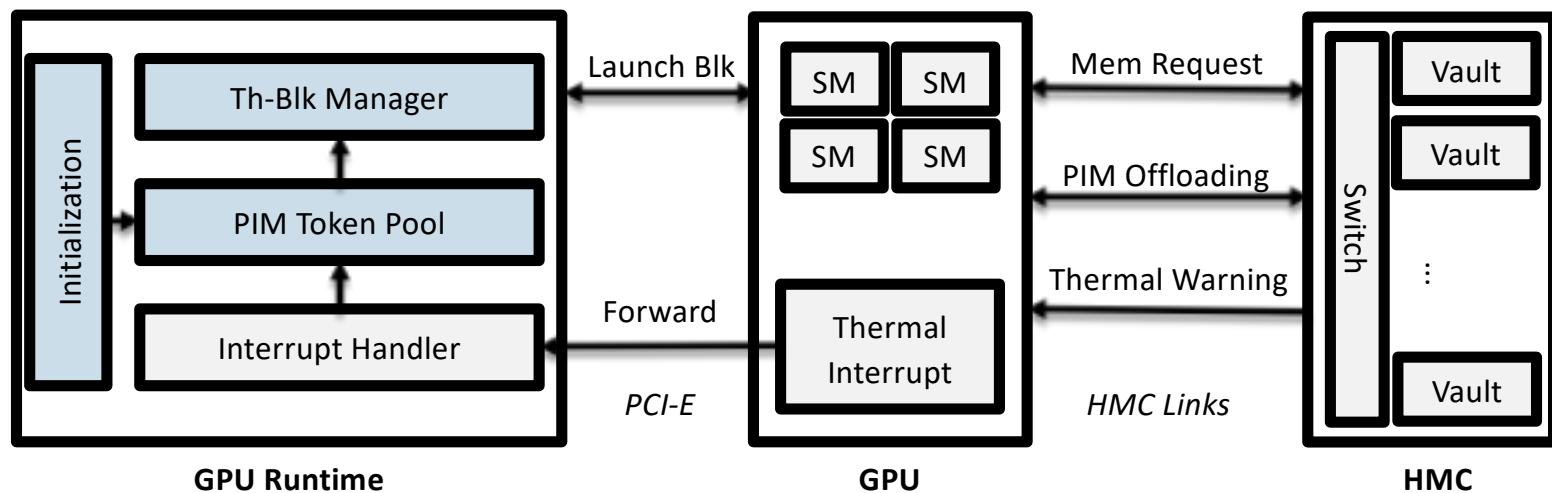


# Software-Based Throttling

14/24

GPU runtime implements some components to control PIM intensity

- PIM Token Pool (PTP)
  - ▶ # of maximum thread blocks that are allowed to use PIM functionality
- Thread Block Manager
  - ▶ Check PTP and launch PIM code if tokens are available
- Initialization
  - ▶ Estimate the initial PTP size based on static analysis at compile time



# Code Generation for non-PIM code

15/24

The GPU compiler generates PIM-enabled and non-PIM kernels at compile time

- Source-to-source translation
- IR-to-IR translation

```
Void cuda_kernel(arg_list)
{
    for (int i=0; i<end; i++)
    {
        uint addr = addrArray[i];
        PIM_Add(addr, 1);
    }
}
```

Original PIM Code

```
void cuda_kernel_np(arg_list)
{
    for (int i=0; i<end; i++)
    {
        uint addr = addrArray[i];
        cuda atomicAdd(addr, 1);
    }
}
```

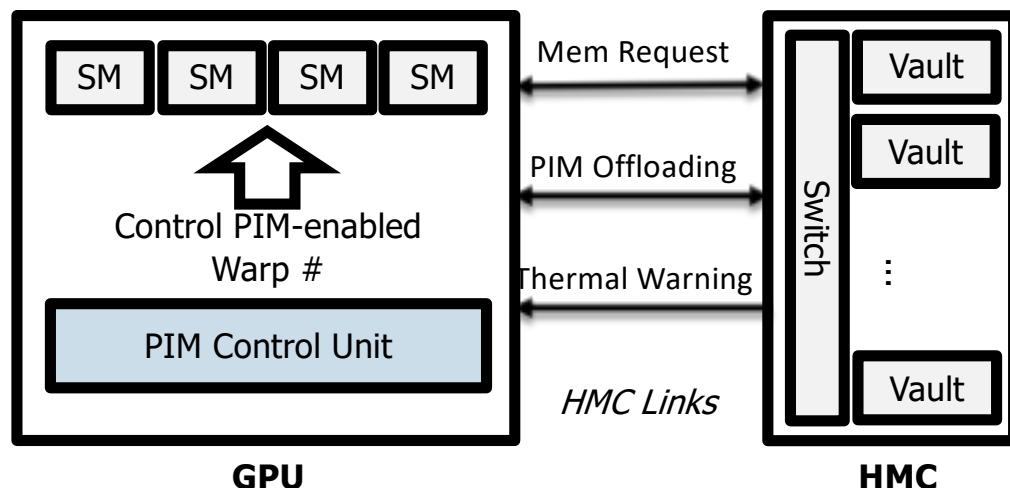
Shadow Non-PIM Code

# Hardware-Based Throttling

16/24

## PIM Control Unit

- Controls # of PIM-enabled warps
- Performs dynamic binary translation
- See the paper for detail!



Type	PIM Instruction	Non-PIM
Arithmetic	Signed Add	atomicAdd
Bitwise	Swap, bit write	atomicExch
Boolean	AND, OR	atomicAND/OR
Comparison	CAS-equal/greater	atomicCAS/Max

## Evaluation

# Methodology

18/24

## Thermal Evaluation

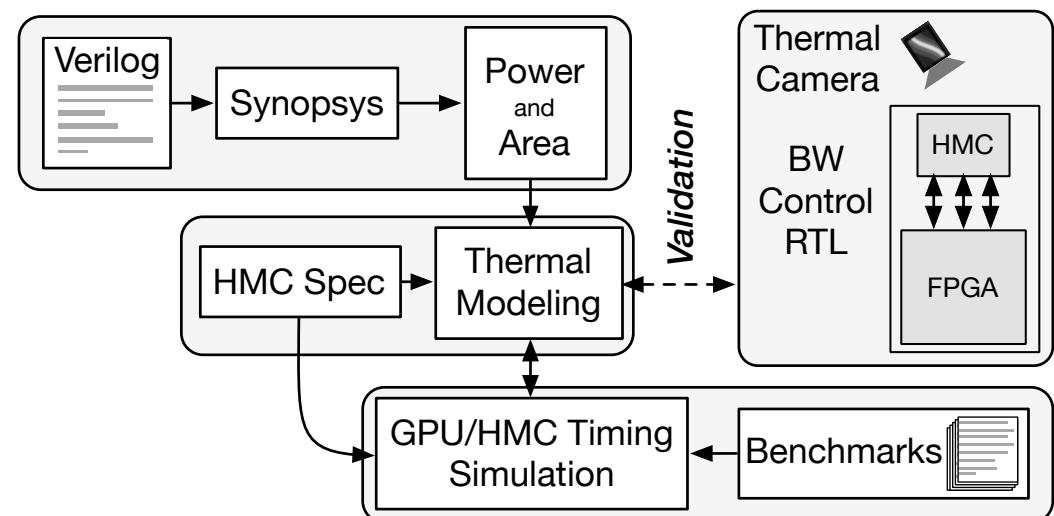
- Temp Measurement: Real HMC 1.1 Platform
- Thermal Modeling: HMC 2.0 using 3D-ICE
- Power & Area: Synopsys (28nm/50nm CMOS)

## Performance Evaluation

- MacSim w/ VaultSim

## Benchmark

- GraphBIG benchmark with LDBC dataset
  - ▶ BFS, SSSP, PageRank, etc...

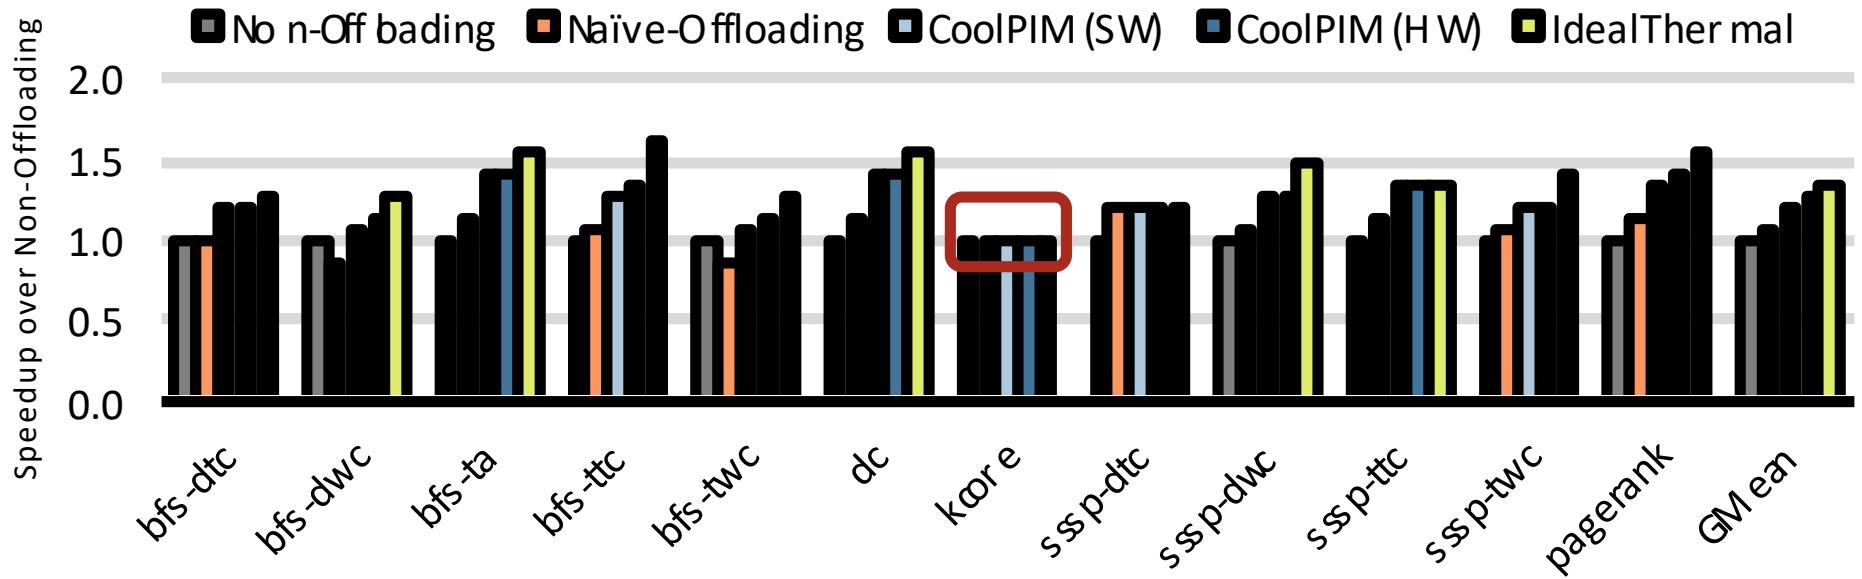


# Performance

19/24

## Speedup over baseline (Non-Offloading)

- **Naïve/SW/HW:** using a commodity-server active heat sync
- **Ideal Thermal:** with unlimited cooling



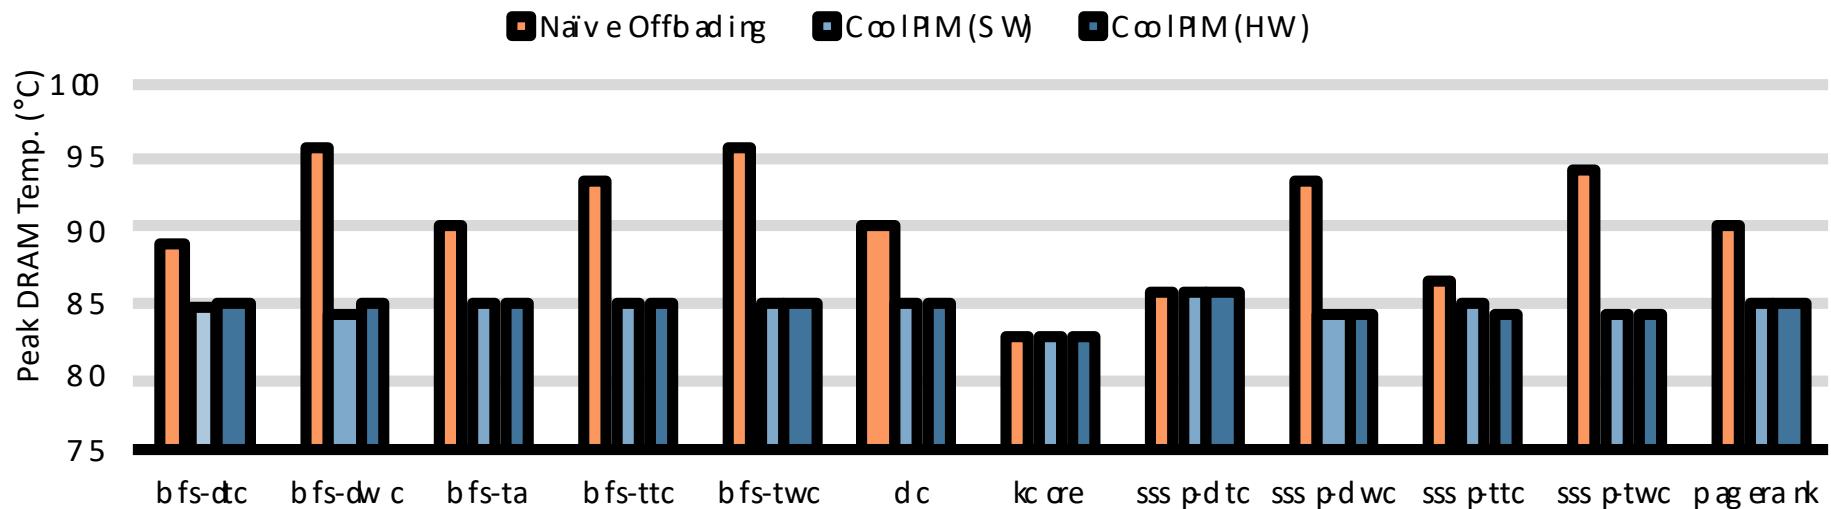
On average, CoolPIM (SW/HW) improves performance by 1.21x/1.25x!

# Thermal Analysis

20/24

## PIM Offloading Rate

- Naïve: 3~4 op/ns → Temperature goes beyond the normal operating region.
- CoolPIM: 1.3 op/ns → No memory performance slowdown



CoolPIM maintains peak DRAM temperature within normal operating temp!

## Conclusion

# Conclusion

---

22/24

**Observation:** PIM integration requires careful thermal consideration

- Naive PIM offloading may cause a thermal issue and degrades overall system performance

**CoolPIM:** Source throttling techniques to control PIM intensity

- Keeps HMC "Cool" to avoid thermal-triggered memory performance degradation

**Results:** CoolPIM improves performance by 1.37x over naïve offloading

- 1.2x over non-offloading on average

# Thank You

# Backup

# Typical Cooling Types

25/24

Type	Thermal Resistance	Cooling Power*
Passive heat sink	4.0 °C/W	0
Low-end active heat sink	2.0 °C/W	1x
Commodity-server active heat sink	0.5 °C/W	104x
High-end heat sink	0.2 °C/W	380x

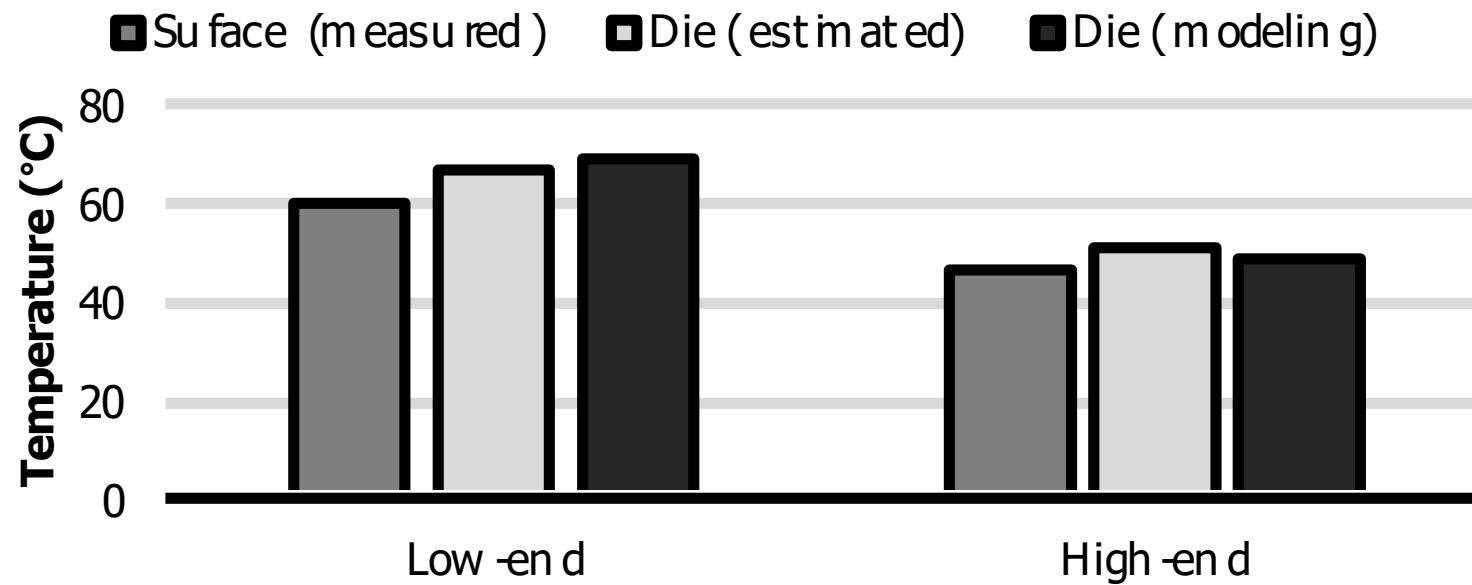
\* We assume the same plate-fin heat sink model for all configurations.

# Thermal Model Validation

26/24

| Validate our thermal evaluation environment

- Model HMC 1.1 temperature and compare with measurements



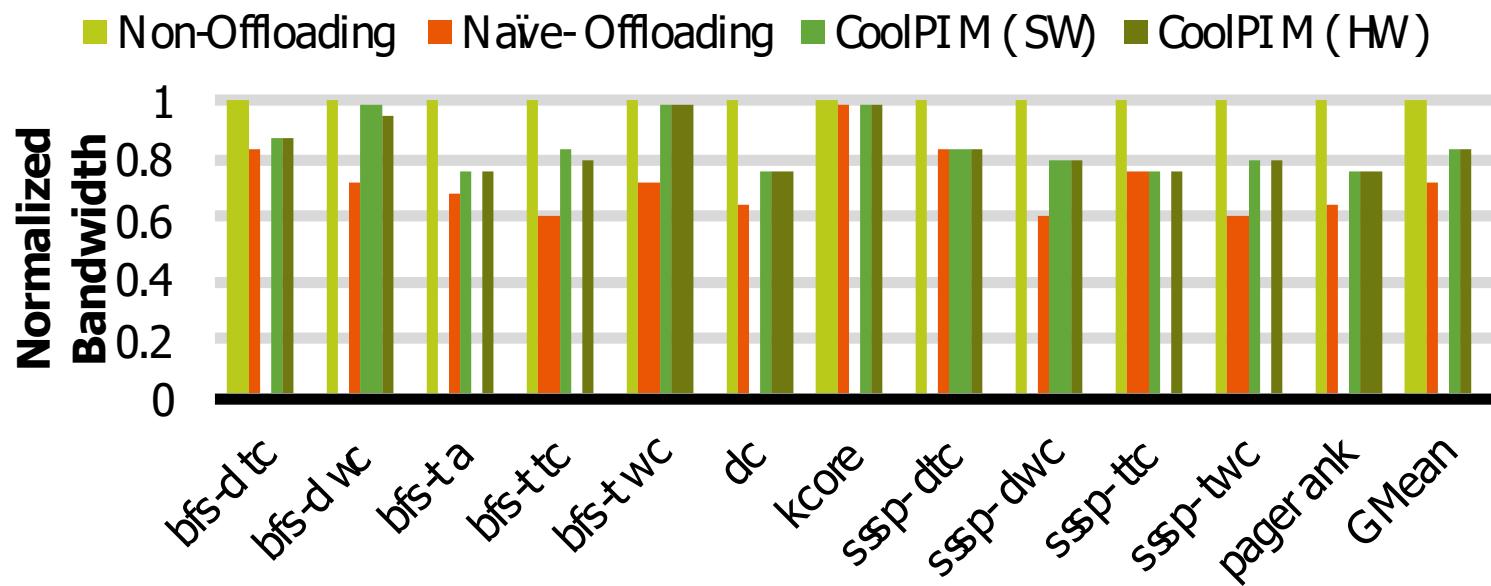
# Evaluation Configuration

27/24

Component	Configuration
Host	GPU, 16 PTX SMs, 32 threads/warp, 1.4GHz 16KB private L1D and 1MB 16-way L2 cache
HMC	8 GB cube, 1 logic die, 8 DRAM dies 32 vaults, 512 DRAM banks $t_{CL}=t_{RCD}=t_{RP}=13.75\text{ns}$ , $t_{RAS}=27.5\text{ns}$ 4 links per package, 120 GB/s per link 80 GB/s data bandwidth per link DRAM Temp. phase: 0-85 °C, 85-95 °C, 95-105 °C 20% DRAM freq reduction (high temp. phases)

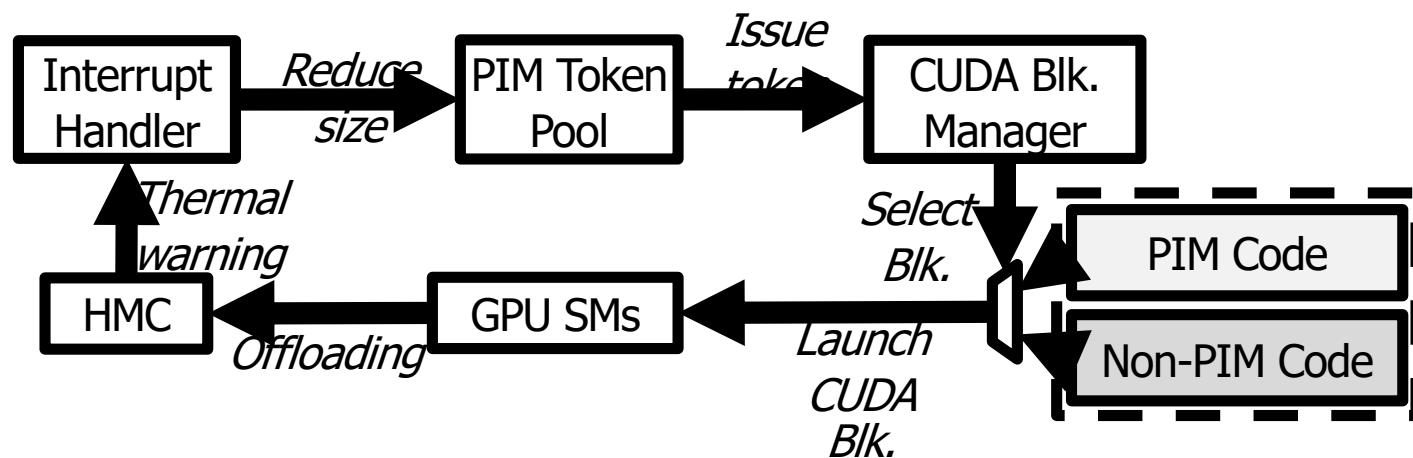
# Bandwidth Consumption

| Bandwidth consumption normalized to baseline (non-offloading)



# Software-Based Throttling

29/24



# Hardware vs Software

Type	Software-Based	HW-Based
Control Granularity	Thread Blocks	Warp
Control Delay	Long Delay	Short Delay
Design Complexity	Low	High

# Hardware vs Software

31/24

